



S05: High Performance Computing with CUDA

Case Study:
Wave propagation through earth

Bernard Deschizeaux
CGGVeritas Oil&Gas service company

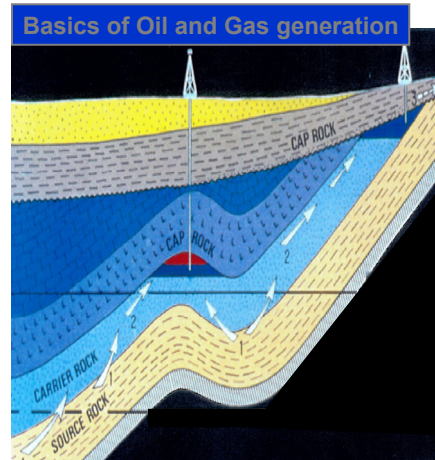
Outline

- What Oil & Gas industry look for
 - Introduction to oil prospection.
 - Seismic prospecting overview.
- What CGGVeritas offers
 - Data acquisition.
 - Data processing.
- Overview of main processing steps
 - Description of the problem to solve.
 - How do we solve it.
- Converting the algorithm in CUDA
 - Kernel example.
- Industrial environment in processing datacenter
 - Using CUDA in hybrid parallel environment.
 - Monitoring tools.
 - GPU server solution.
- Conclusion

What Oil & Gas look for



- Oil is generated deep down the earth.
- Oil migrate up through porous rocks.
- Oil is trapped by none porous rocks with specific structural shapes.
- To find oil the knowledge of the earth structure and rocks property is mandatory.




S05: High Performance Computing with CUDA



Economical pressure for accuracy



- Mature fields need from 10 to 100 wells.
- Typical cost for drilling a well is 10 millions dollars.
- At actual stat of the art 20% of the drilled wells are lost!
- Wells need to be positioned at a few meters error, below 1km of sea water and few km of rocks.





S05: High Performance Computing with CUDA




Earth studies: Seismic survey



- Wave propagation in the earth is related to rocks properties



- Arrival time of reflection gives structural information.
- Amplitude of reflection gives information on rocks properties. (wave velocity, rock density)



S05: High Performance Computing with CUDA

5



CGGVeritas Today




S05: High Performance Computing with CUDA

6



1 seismic survey = 10TB of data



- One shot ~ 5000 records of 10KB.
- Typical acquisition: 10^5 shots to sample the earth every 25m
- 10TB of raw data.



S05: High Performance Computing with CUDA

Seismic processing




- 10TB of raw data.
- Every byte of data will need 10^6 operations.
- Through the processing, data size will be increased tenfold.
- CGGV processing capability is centralized in main datacenters providing hundreds of Teraflops of computer capability and PetaBytes of storage.




S05: High Performance Computing with CUDA

Converting arrival time into depth



- In case of radar distances are simple to extract from arrival time: $D = T * V/2$.

- Constant velocity in air.
- Straight line propagation.






- In Seismic exploration the problem is far more complex.

- Far more complex reflections.
- Velocity field is unknown.
- Velocity field is inhomogeneous.
⇒ Ray path are not straight.

S05: High Performance Computing with CUDA

9 

Migration: dips & wave velocity effect




- For correct image we need the velocity field.
- The wave needs to be propagated in all directions.

S05: High Performance Computing with CUDA

10 


Migration results example



S05: High Performance Computing with CUDA

11 SC07

The wave propagation algorithm



- Finite difference algorithm used in frequency to propagation from one depth to the next.

- Propagator expressed in the spatial direction using Chebychev polynomials of second derivative.

$$G(\Delta) = \sum_{n=0}^{n=N} a_{\omega/V}(n) T_n(\Delta)$$

- $T_n(x)$: Chebychev polynomial of degree n computed using the recursive equation:

$$T_n(x) = 2x.T_{n-1}(x) - T_{n-2}(x)$$

- Δ the second derivative is approximated by a 2D cross filter.

Reference: Soubaras, R., 1996. "Explicit 3-D migration using equiripple polynomial expansion and Laplacian synthesis". Geophysics, Volume 61, pp. 1386_1393

S05: High Performance Computing with CUDA

12 SC07


Propagator algorithm



```

T(x,y) = Wavez(x,y);
TT(x,y) = Δ ⊗ Wavez(x,y);
Wavezdz(x,y) = α1 * T(x,y)
                + α2 * TT(x,y);

for (n = 2; n < NMAX; n++){
    // Compute the Chebychev polynomial TTT
    TTT(x,y) = 2 * Δ ⊗ TT(x,y) - T(x,y);
    // Store Chebychev results for next iteration.
    T(x,y) = TT(x,y);
    TT(x,y) = TTT(x,y);
    // update result.
    Wavezdz(x,y) += αn * TTT(x,y);
};
    
```




S05: High Performance Computing with CUDA

13



10¹¹ propagations for the final image



S05: High Performance Computing with CUDA

14



CUDA propagator kernel



Block geometry: $64 \times 8 = 512$ threads


- **Load:**

- Each thread reads one value in each quarter of the block data.

- **Processing:**

- Each thread applies the cross operator to one value in each third of the block data.

Random access for the kernel.
But sequential access for all warps of a block.



S05: High Performance Computing with CUDA

15 SC07

Complex performance analyses



- Kernels speedup varying from none to one order of magnitude.
- More than 100 parameters.
 - Large varieties of survey geometry.
 - Input frequency plan size from 50×50 to 500×500 .
 - Pre and Post processing percentage (Amdahl law).
 - Interpolation increment, operator length
- Global speedup even more difficult to extract.
 - Finding representative dataset.
 - Finding representative parameters (empirical).
 - Knowing that even 2X speedup will change drastically the whole methodology.

S05: High Performance Computing with CUDA

16 SC07

CUDA development



- Development is not straightforward.
- Programming model very powerful and flexible.
- Lack of debugger and profiler.
- Some weakness in error handling.
- CUDA still in the evolution curve (more functionality and less constraint yet to come).
- Preliminary performance results on individual kernels very encouraging.
- Ongoing benchmark on different datasets and with realistic parameters.
- Hardware GPU server solution still in evaluation.

S05: High Performance Computing with CUDA

17



Development for processing hub



The prototype is included in a composite multi-layers parallelism

- Many processing project running concurrently.
- Large processing step split and schedule on different PC cluster racks.
- MPI used to parallelize on nodes in a rack.
- OpenMP used to parallelize on cores in a node and to schedule multi GPU.
- CUDA used to parallelize on GPU.

S05: High Performance Computing with CUDA


18



Monitoring 1000th of Nodes



- For a processing hub of many thousands of nodes, hardware failures occur every day.
- Monitoring tools are needed to detect and analyze problems and efficiency.
- We cannot afford to lose more than a few hours processing because of failure.
- Failure rate of every piece of hardware needed to scale job size and backup save rate.



S05: High Performance Computing with CUDA

19



GPU in a processing datacenter



- Performance measurement
 - Speed Up.
 - Overall percentage of applications using GPU.
- Large datacenter constraints.
 - Space.
 - Power supply.
 - Cooling system.
 - Maintenance.
 - Price.
- What GPU solution need to provide
 - Compact (1u server).
 - Reasonable power needs.
 - Integrated in the general cooling system.
 - Stable for sustain use.
 - Inexpensive.

S05: High Performance Computing with CUDA

20



Conclusion



- Collaboration with NVIDIA is fruitful.
- Industrial time consuming algorithm successfully prototype with CUDA.
- Extensive benchmark ongoing on real scale datasets.
- Design for GPU server still in progress.
- Capability trend of GPU versus CPU show that this is a promising technology for number crunching solutions.

